

CNN Experiments on CIFAR-10

Devansh Sharma

November 18, 2025

1. Introduction

This project evaluates how increasing CNN depth affects classification accuracy on the CIFAR-10 dataset. Three different convolutional neural networks were implemented:

- **CNN_Small** – 3 convolutional layers + 2 fully connected layers
- **CNN_Large** – 6 convolutional layers + 2 fully connected layers
- **CNN_XL** – 9 convolutional layers with BatchNorm + 2 fully connected layers

All experiments used:

- Cross-entropy loss
- SGD optimizer ($lr=0.01$, $momentum=0.9$, $weight\ decay=5e-4$)
- StepLR scheduler
- 30 training epochs

Performance was tracked using Weights & Biases (W&B).

2. Model Architectures

2.1 CNN_Small

- $\text{Conv}(3 \rightarrow 32)$, $\text{Conv}(32 \rightarrow 32)$, MaxPool
- $\text{Conv}(32 \rightarrow 64)$, $\text{Conv}(64 \rightarrow 64)$, MaxPool
- Fully-connected: $1024 \rightarrow 256 \rightarrow 10$

2.2 CNN_Large

- Same structure as Small, but an additional block:
- $\text{Conv}(64 \rightarrow 128)$, $\text{Conv}(128 \rightarrow 128)$, MaxPool
- Fully-connected: $2048 \rightarrow 256 \rightarrow 10$

2.3 CNN_XL

- All blocks upgraded with BatchNorm
- 64-channel, 128-channel, and 256-channel double conv blocks
- Fully-connected: $4096 \rightarrow 512 \rightarrow 10$

3. Results

Final test accuracy after 30 epochs:

Model	Train Accuracy	Test Accuracy
CNN_Small	0.8183	0.8050
CNN_Large	0.8291	0.8216
CNN_XL	0.9412	0.8942

Table 1: Accuracy comparison across CNN model depths.

Observations:

- Adding depth improves performance (Small \rightarrow Large).
- Adding depth + BatchNorm dramatically improves accuracy (XL model).
- CNN_XL shows both the highest training stability and lowest loss curves.

4. Training Curves

4.1 Accuracy Curves

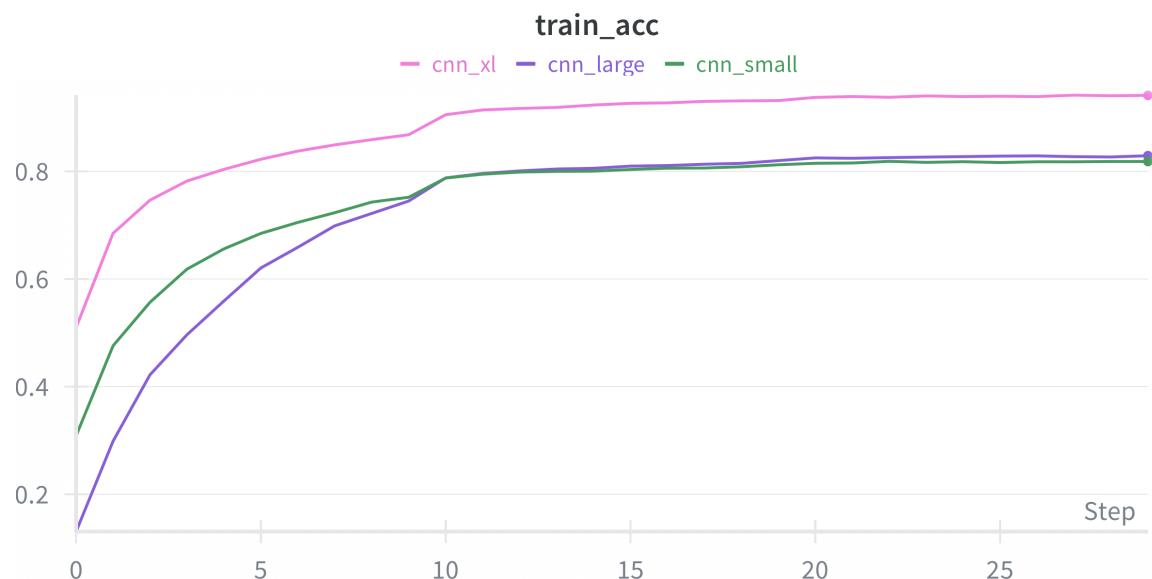


Figure 1: Training Accuracy

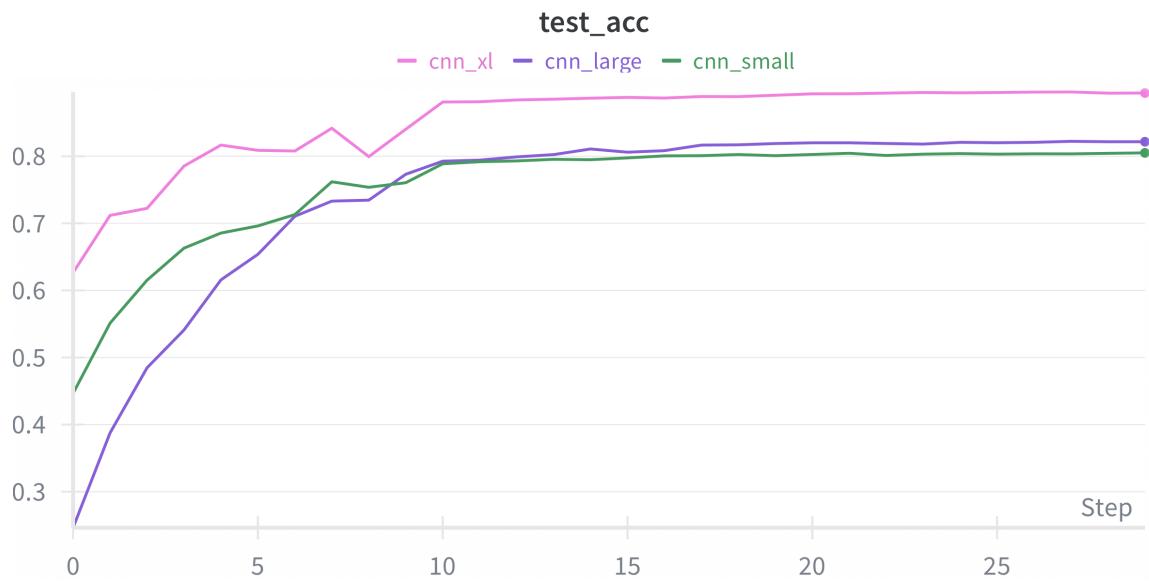


Figure 2: Test Accuracy

4.2 Loss Curves



Figure 3: Training Loss

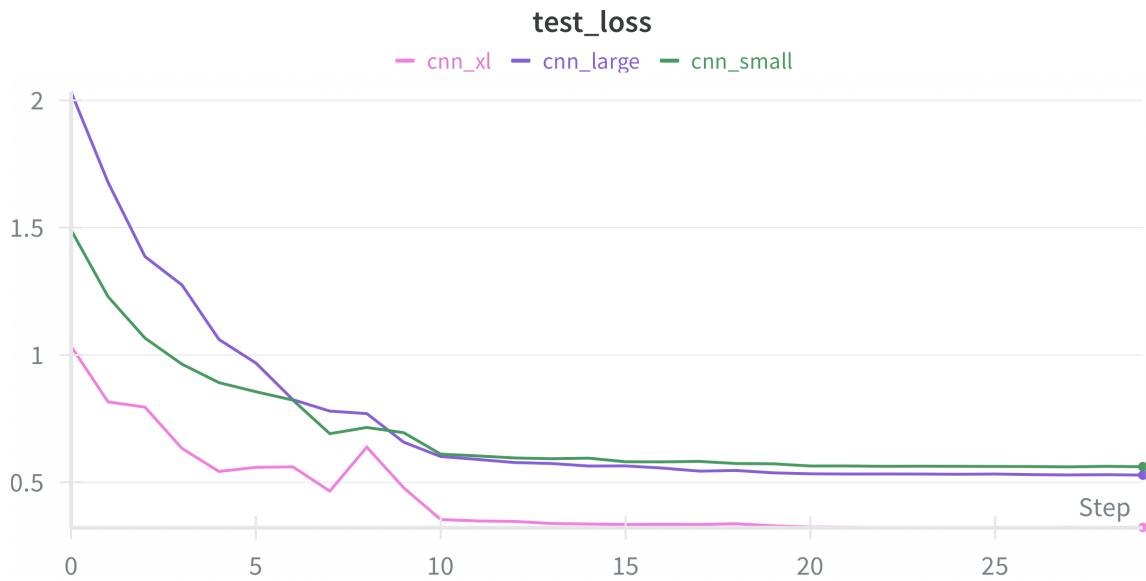


Figure 4: Test Loss

5. Conclusion

The experiments clearly show that increasing CNN depth improves accuracy on CIFAR-10. However, the largest improvement came from adding Batch Normalization to each block in the CNN_XL model. This resulted in:

- Faster convergence
- Higher final accuracy
- More stable loss curves

CNN_XL achieved a strong **89.4%** test accuracy in only 30 epochs, demonstrating the importance of deeper and properly normalized architectures.